

Hyperthyroidism Detection: A Comparative Analysis of Classifiers

This project aims to develop a predictive model for diagnosing hyperthyroidism in patients using machine learning techniques. We will utilize a dataset containing various features related to thyroid function tests, symptoms, and medical history to train and evaluate several classifiers. The study will focus on comparing the performance of classifiers such as Multinomial Naive Bayes, Support Vector Machine (SVM), Random Forest, and Multilayer Perceptron (MLP) to determine the most effective model for accurate hyperthyroidism prediction. The project's ultimate goal is to contribute to improved medical diagnostics and decision-making in the field of thyroid disorders.

Import Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import MultiLabelBinarizer
from sklearn.multiclass import OneVsRestClassifier
import warnings
warnings.filterwarnings('ignore')
```

Uploading CSV


uploading csv, storing in dataframe and displaying first 5 entries of dataframe

```
In [2]: df = pd.read_csv('thyroidDF.csv')
df.head()
```

Out[2]:

	age	sex	on_thyroxine	query_on_thyroxine	on_antithyroid_meds	sick	pregnant	thyroid
0	29	F	f	f	f	f	f	f
1	29	F	f	f	f	f	f	f
2	41	F	f	f	f	f	f	f
3	36	F	f	f	f	f	f	f
4	32	F	f	f	f	f	f	f

5 rows × 31 columns



.shape

returns a tuple representing the dimensions (number of rows and columns) of a pandas DataFrame.

```
In [3]: df.shape
```

Out[3]: (9172, 31)

features types

assing and displaying features types.

```
In [4]: feature_types = {
    'age': 'Numerical',
    'sex': 'Nominal Categorical',
    'on_thyroxine': 'Boolean',
    'query on thyroxine': 'Boolean',
    'on antithyroid meds': 'Boolean',
    'sick': 'Boolean',
    'pregnant': 'Boolean',
    'thyroid_surgery': 'Boolean',
    'I131_treatment': 'Boolean',
    'query_hypothyroid': 'Boolean',
    'query_hyperthyroid': 'Boolean',
    'lithium': 'Boolean',
    'goitre': 'Boolean',
    'tumor': 'Boolean',
    'hypopituitary': 'Numerical',
    'psych': 'Boolean',
    'TSH_measured': 'Boolean',
    'TSH': 'Numerical',
    'T3_measured': 'Boolean',
    'T3': 'Numerical',
    'TT4_measured': 'Boolean',
    'TT4': 'Numerical',
    'T4U_measured': 'Boolean',
    'T4U': 'Numerical',
    'FTI_measured': 'Boolean',
    'FTI': 'Numerical',
    'TBG_measured': 'Boolean',
    'TBG': 'Numerical',
    'referral_source': 'Nominal Categorical',
    'target': 'Nominal Categorical',
    'patient_id': 'Nominal Categorical'
}

# Display feature types
for column in df.columns:
    if column in feature_types:
        print(f"{column}: {feature_types[column]}")
    else:
        print(f"{column}: Not specified")
```

```

age: Numerical
sex: Nominal Categorical
on_thyroxine: Boolean
query_on_thyroxine: Not specified
on_antithyroid_meds: Not specified
sick: Boolean
pregnant: Boolean
thyroid_surgery: Boolean
I131_treatment: Boolean
query_hypothyroid: Boolean
query_hyperthyroid: Boolean
lithium: Boolean
goitre: Boolean
tumor: Boolean
hypopituitary: Numerical
psych: Boolean
TSH_measured: Boolean
TSH: Numerical
T3_measured: Boolean
T3: Numerical
TT4_measured: Boolean
TT4: Numerical
T4U_measured: Boolean
T4U: Numerical
FTI_measured: Boolean
FTI: Numerical
TBG_measured: Boolean
TBG: Numerical
referral_source: Nominal Categorical
target: Nominal Categorical
patient_id: Nominal Categorical

```

dropping


dropping column patient_id.

```
In [5]: df.drop(['patient_id'], axis = 1, inplace = True)
df.head()
```

Out[5]:

	age	sex	on_thyroxine	query_on_thyroxine	on_antithyroid_meds	sick	pregnant	thyroic
0	29	F	f	f	f	f	f	f
1	29	F	f	f	f	f	f	f
2	41	F	f	f	f	f	f	f
3	36	F	f	f	f	f	f	f
4	32	F	f	f	f	f	f	f

5 rows × 30 columns



.info

provides a concise summary of information about a pandas DataFrame, including data types, non-null counts, and memory usage.

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9172 entries, 0 to 9171
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                   9172 non-null   int64
1   sex                                   8865 non-null   object
2   on_thyroxine                         9172 non-null   object
3   query_on_thyroxine                  9172 non-null   object
4   on_antithyroid_meds                 9172 non-null   object
5   sick                                 9172 non-null   object
6   pregnant                             9172 non-null   object
7   thyroid_surgery                     9172 non-null   object
8   I131_treatment                      9172 non-null   object
9   query_hypothyroid                   9172 non-null   object
10  query_hyperthyroid                   9172 non-null   object
11  lithium                              9172 non-null   object
12  goitre                               9172 non-null   object
13  tumor                                9172 non-null   object
14  hypopituitary                       9172 non-null   object
15  psych                                9172 non-null   object
16  TSH_measured                        9172 non-null   object
17  TSH                                  8330 non-null   float64
18  T3_measured                         9172 non-null   object
19  T3                                   6568 non-null   float64
20  TT4_measured                        9172 non-null   object
21  TT4                                  8730 non-null   float64
22  T4U_measured                        9172 non-null   object
23  T4U                                  8363 non-null   float64
24  FTI_measured                        9172 non-null   object
25  FTI                                  8370 non-null   float64
26  TBG_measured                        9172 non-null   object
27  TBG                                  349 non-null    float64
28  referral_source                     9172 non-null   object
29  target                              9172 non-null   object
dtypes: float64(6), int64(1), object(23)
memory usage: 2.1+ MB
```

missing data analysis and handling

calculates and returns the count of missing (null) values for each column in a DataFrame, allowing for easy identification of data gaps.

In [7]: `df.isnull().sum()`

```
Out[7]: age                0
sex              307
on_thyroxine     0
query_on_thyroxine  0
on_antithyroid_meds  0
sick             0
pregnant        0
thyroid_surgery  0
I131_treatment   0
query_hypothyroid  0
query_hyperthyroid  0
lithium         0
goitre          0
tumor           0
hypopituitary   0
psych           0
TSH_measured    0
TSH             842
T3_measured     0
T3             2604
TT4_measured    0
TT4            442
T4U_measured    0
T4U            809
FTI_measured    0
FTI            802
TBG_measured    0
TBG            8823
referral_source  0
target          0
dtype: int64
```

```
In [8]: def fill_column(column_A, column_B):
    if column_A == 'f':
        return 0.0
    else:
        return column_B
```

```
In [9]: df['TSH'] = df.apply(lambda row: fill_column(row['TSH_measured'], row['TSH']),
df['T3'] = df.apply(lambda row: fill_column(row['T3_measured'], row['T3']),
df['TT4'] = df.apply(lambda row: fill_column(row['TT4_measured'], row['TT4']),
df['T4U'] = df.apply(lambda row: fill_column(row['T4U_measured'], row['T4U']),
df['FTI'] = df.apply(lambda row: fill_column(row['FTI_measured'], row['FTI']),
df['TBG'] = df.apply(lambda row: fill_column(row['TBG_measured'], row['TBG']),
```

```
In [10]: df = df.fillna("N")
```

dupliacted data analysis and handling

calculates and returns the count of missing (null) values for each column in a DataFrame , helping to assess the impact of data removal on missing data patterns.

```
In [11]: df.duplicated().sum()
```

```
Out[11]: 10
```


```
In [12]: df = df.drop_duplicates()
```

```
In [13]: df_visualization = df.copy()
```

Visualizations

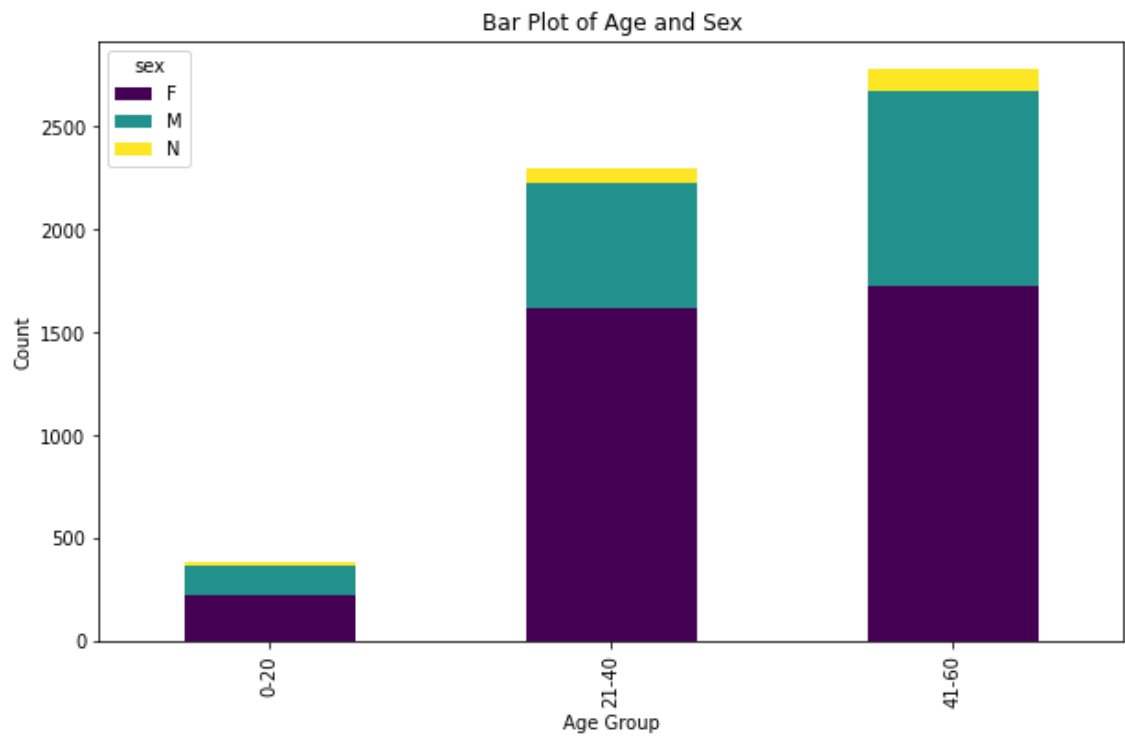
Visualizing Gender Distribution

```
In [14]: type_counts = df_visualization['sex'].value_counts()  
plt.figure(figsize=(6, 6))  
plt.pie(type_counts, labels=type_counts.index, autopct='%1.1f%%', startangle=0)  
plt.title('Percentage Distribution of Sex')  
plt.axis('equal')  
plt.show()
```



Exploring Age and Gender Distribution

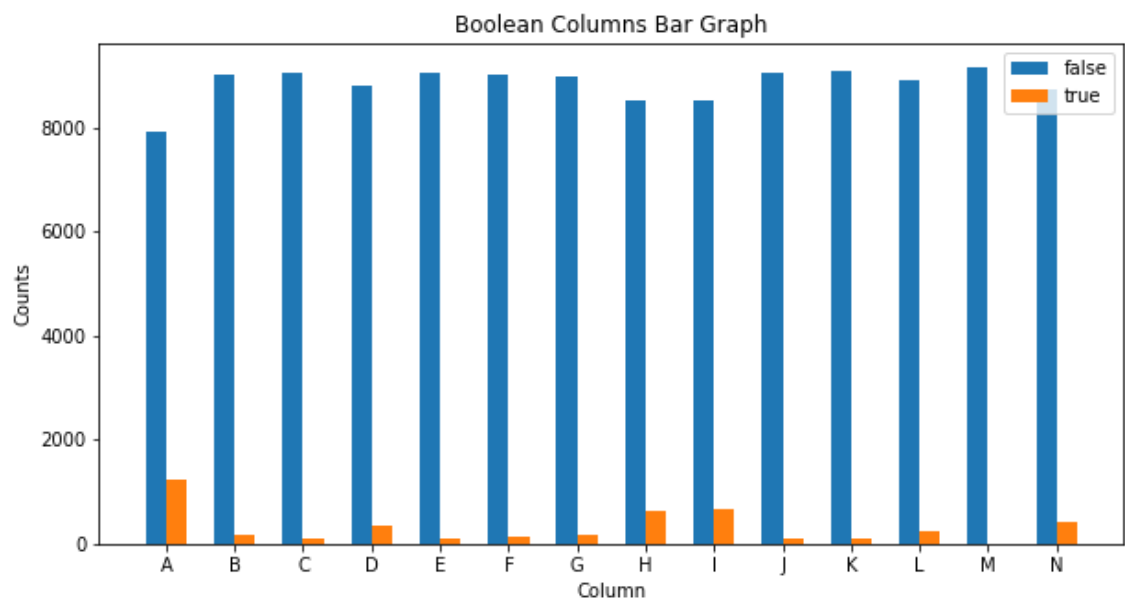
```
In [15]: bins = [0, 20, 40, 60]
labels = ['0-20', '21-40', '41-60']
df_visualization['age_group'] = pd.cut(df_visualization['age'], bins=bins,
grouped_data = df_visualization.groupby(['age_group', 'sex']).size().unstack()
grouped_data.plot(kind='bar', stacked=True, colormap='viridis', figsize=(10, 6))
plt.xlabel('Age Group')
plt.ylabel('Count')
plt.title('Bar Plot of Age and Sex')
plt.show()
df_visualization.drop(['age_group'], axis = 1, inplace = True)
```



Analyzing Boolean Column Distribution

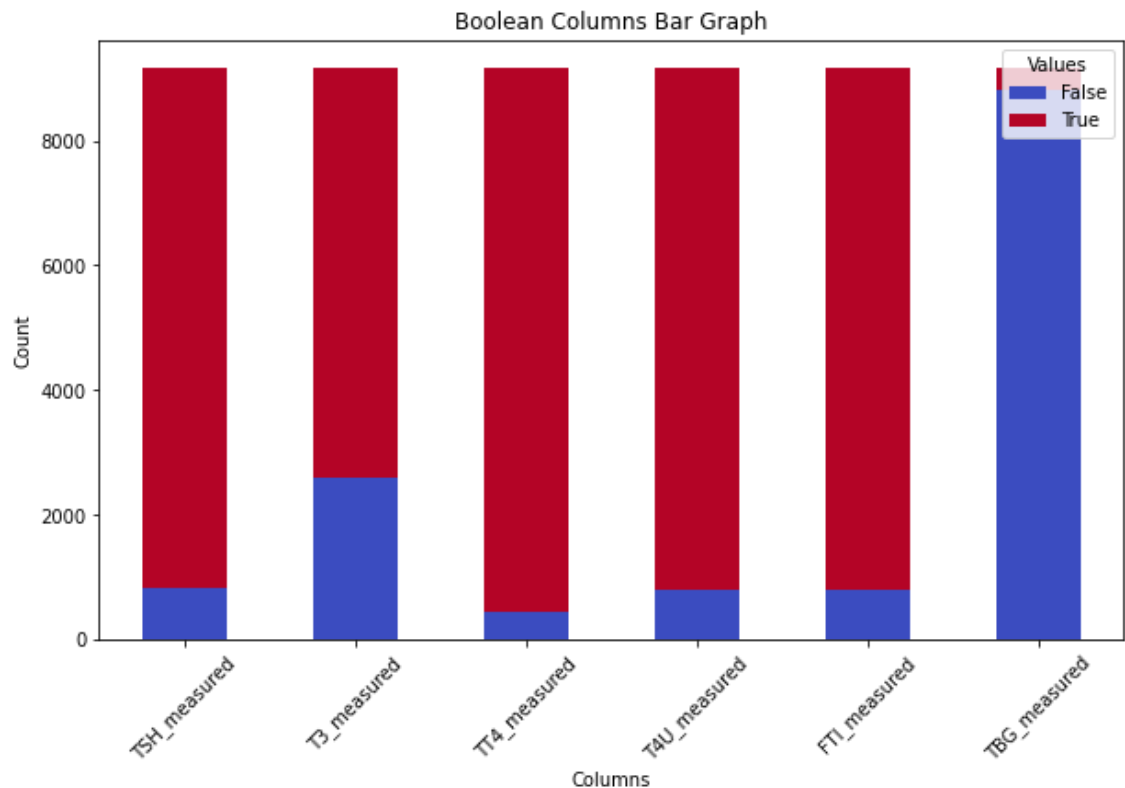
```
In [16]: ▶ selected_columns = ['on_thyroxine', 'query_on_thyroxine', 'on_antithyroid_m',
                              'thyroid_surgery', 'I131_treatment', 'query_hypothyroid', 'hypopituitary', 'psych']

counts = df_visualization[selected_columns].apply(pd.Series.value_counts).T
N = 14
blue_bar = (counts['f'])
orange_bar = (counts['t'])
ind = np.arange(N)
plt.figure(figsize=(10,5))
width = 0.3
plt.bar(ind, blue_bar, width, label='false')
plt.bar(ind + width, orange_bar, width, label='true')
plt.xlabel('Column')
plt.ylabel('Counts')
plt.title('Boolean Columns Bar Graph')
plt.xticks(ind + width / 2, ('A', 'B', 'C', 'D', 'E',
                              'F', 'G', 'H', 'I', 'J', 'K', 'L',
                              'M', 'N'))
plt.legend(loc='best')
plt.show()
```



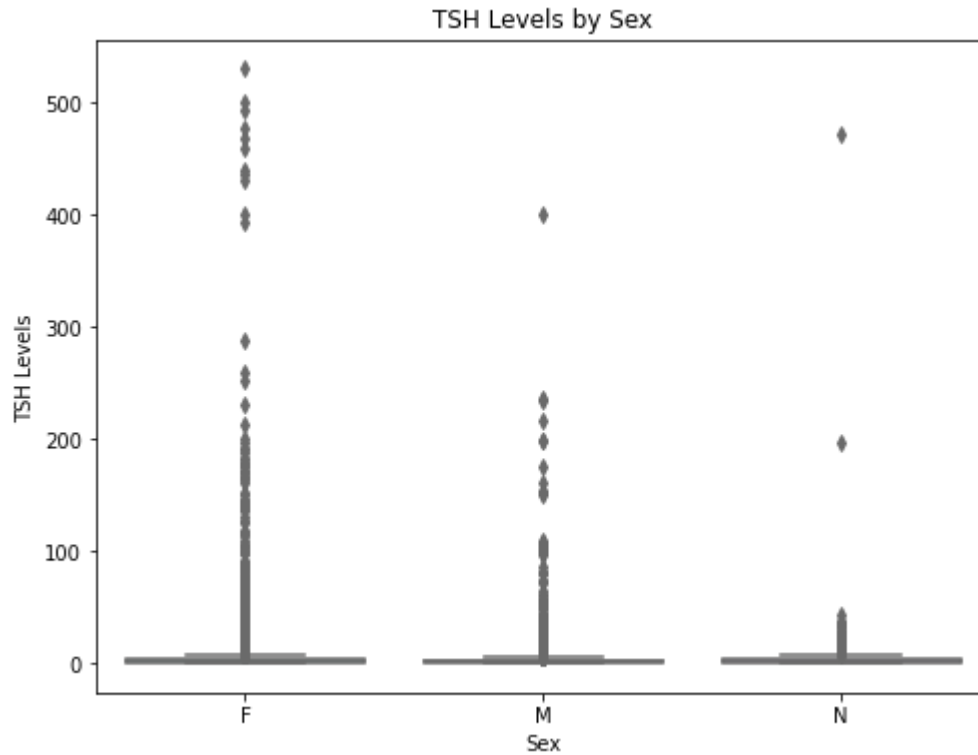
Exploring Measured Thyroid Parameters

```
In [17]: ▶ selected_columns = ['TSH_measured', 'T3_measured', 'TT4_measured', 'T4U_measured', 'FTI_measured', 'TBG_measured']
counts = df_visualization[selected_columns].apply(pd.Series.value_counts).T
counts.plot(kind='bar', stacked=True, colormap='coolwarm', figsize=(10, 6))
plt.title('Boolean Columns Bar Graph')
plt.xlabel('Columns')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.legend(title='Values', labels=['False', 'True'], loc='upper right')
plt.show()
```



Comparing TSH Levels by Gender

```
In [18]: ▶ plt.figure(figsize=(8, 6))
sns.boxplot(x='sex', y='TSH', data=df_visualization, palette='Set3')
plt.xlabel('Sex')
plt.ylabel('TSH Levels')
plt.title('TSH Levels by Sex')
plt.show()
```




```
In [19]: ▶ def process_target(target):
          classes = target.split('|')
          return classes
```

Exploring Top Target Variables Distribution

```
In [20]: ▶ df_visualization['Processed_Target'] = df_visualization['target'].apply(proc
all_classes = [cls for classes in df_visualization['Processed_Target'] for
class_counts = pd.Series(all_classes).value_counts()
top_5_classes = class_counts.head(5)
other_count = class_counts[5:].sum()
top_5_classes['Others'] = other_count
plt.figure(figsize=(10, 6))
plt.pie(top_5_classes, labels=top_5_classes.index, autopct='%1.1f%%', star
plt.title('Percentage of Top 5 Target Variables (with Others)')
plt.axis('equal')
plt.show()
df_visualization.drop(['Processed_Target'], axis = 1, inplace = True)
df_visualization.drop(['referral_source'], axis = 1, inplace = True)
```


Percentage of Top 5 Target Variables (with Others)



Correlation Heatmap of Attributes

```
In [21]: sns.set(rc={'figure.figsize':(10, 10)})  
sns.set(font_scale=1.2)  
sns.heatmap(data=df_visualization.corr(), cmap='Blues', annot=True)
```

Out[21]: <Axes: >



Data Preprocessing

Label Encoding for Categorical Data

```
In [22]: ▶ label_encoder = LabelEncoder()
df_visualization['sex'] = label_encoder.fit_transform(df_visualization['sex'])
df_visualization['on_thyroxine'] = label_encoder.fit_transform(df_visualization['on_thyroxine'])
df_visualization['query_on_thyroxine'] = label_encoder.fit_transform(df_visualization['query_on_thyroxine'])
df_visualization['on_antithyroid_meds'] = label_encoder.fit_transform(df_visualization['on_antithyroid_meds'])
df_visualization['sick'] = label_encoder.fit_transform(df_visualization['sick'])
df_visualization['pregnant'] = label_encoder.fit_transform(df_visualization['pregnant'])
df_visualization['thyroid_surgery'] = label_encoder.fit_transform(df_visualization['thyroid_surgery'])
df_visualization['I131_treatment'] = label_encoder.fit_transform(df_visualization['I131_treatment'])
df_visualization['query_hypothyroid'] = label_encoder.fit_transform(df_visualization['query_hypothyroid'])
df_visualization['query_hyperthyroid'] = label_encoder.fit_transform(df_visualization['query_hyperthyroid'])
df_visualization['lithium'] = label_encoder.fit_transform(df_visualization['lithium'])
df_visualization['goitre'] = label_encoder.fit_transform(df_visualization['goitre'])
df_visualization['tumor'] = label_encoder.fit_transform(df_visualization['tumor'])
df_visualization['hypopituitary'] = label_encoder.fit_transform(df_visualization['hypopituitary'])
df_visualization['psych'] = label_encoder.fit_transform(df_visualization['psych'])
df_visualization['TSH_measured'] = label_encoder.fit_transform(df_visualization['TSH_measured'])
df_visualization['T3_measured'] = label_encoder.fit_transform(df_visualization['T3_measured'])
df_visualization['TT4_measured'] = label_encoder.fit_transform(df_visualization['TT4_measured'])
df_visualization['T4U_measured'] = label_encoder.fit_transform(df_visualization['T4U_measured'])
df_visualization['FTI_measured'] = label_encoder.fit_transform(df_visualization['FTI_measured'])
df_visualization['TBG_measured'] = label_encoder.fit_transform(df_visualization['TBG_measured'])
```

Feature Scaling for Model Enhancement

Min-Max scaling to normalize selected columns within the DataFrame.

```
In [23]: ▶ def column_to_normalize(df, column):
min_value = df[column].min()
max_value = df[column].max()
df[column] = (df[column] - min_value) / (max_value - min_value)
```

```
In [24]: ▶ column_to_normalize(df_visualization, 'age')
column_to_normalize(df_visualization, 'TSH')
column_to_normalize(df_visualization, 'T3')
column_to_normalize(df_visualization, 'TT4')
column_to_normalize(df_visualization, 'T4U')
column_to_normalize(df_visualization, 'FTI')
column_to_normalize(df_visualization, 'TBG')
```

Multi-Label Hyperthyroidism Classification

the 'target' column in the DataFrame 'df_visualization' for multi-label classification by replacing hyphens with 'X' and splitting the labels separated by '|' using MultiLabelBinarizer.

```
In [25]: df_visualization['target'] = df_visualization['target'].replace('-', 'X')
df_visualization['target'] = df_visualization['target'].apply(lambda x: x.split(',')
mlb = MultiLabelBinarizer()
```

Model Preparing

Feature-Target Split

```
In [26]: y = mlb.fit_transform(df_visualization['target'])
X = df_visualization.drop(['target'], axis=1)
X.columns = X.columns.astype(str)
```

Train-Test Split

```
In [27]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Model Training and Testing


Support Vector Machine (SVM) Classifier and Accuracy Visualization

A Support Vector Machine (SVM) classifier is trained and evaluated on the provided training and testing datasets (X_train, y_train, X_test, y_test). The classifier's accuracy on both the testing and training sets is calculated and displayed.

```
In [44]: ► svm_classifier = OneVsRestClassifier(SVC(kernel='linear'))
svm_classifier.fit(X_train, y_train)
svm_preds = svm_classifier.predict(X_test)
svm_accuracy = accuracy_score(y_test, svm_preds)
svm_train_preds = svm_classifier.predict(X_train)
svm_train_accuracy = accuracy_score(y_train, svm_train_preds)
print(f"SVM Testing Accuracy: {svm_accuracy}")
print(f"SVM Training Accuracy: {svm_train_accuracy}")
categories = ['Testing Accuracy', 'Training Accuracy']
values = [svm_accuracy, svm_train_accuracy]
plt.figure(figsize=(8, 6))
sns.barplot(x=categories, y=values, palette='pastel')
plt.ylabel('Accuracy')
plt.title('SVM Classifier Accuracy')
plt.ylim([0, 1]) # Set y-axis limits to match accuracy range (0 to 1)
plt.show()
```

SVM Testing Accuracy: 0.7463175122749591

SVM Training Accuracy: 0.7553554373038613




Random Forest Classification and Accuracy Visualization

A Random Forest classifier to the training data, calculates both testing and training accuracies, and then displays the results. It also visualizes the accuracies using a bar chart, illustrating the performance of the Random Forest model on the dataset.


```
In [45]: rf_classifier = OneVsRestClassifier(RandomForestClassifier(max_depth=5))
rf_classifier.fit(X_train, y_train)
rf_preds = rf_classifier.predict(X_test)
rf_accuracy = accuracy_score(y_test, rf_preds)
rf_train_preds = rf_classifier.predict(X_train)
rf_train_accuracy = accuracy_score(y_train, rf_train_preds)
print(f"RF Testing Accuracy: {rf_accuracy}")
print(f"RF Training Accuracy: {rf_train_accuracy}")
categories = ['Testing Accuracy', 'Training Accuracy']
values = [rf_accuracy, rf_train_accuracy]
plt.figure(figsize=(8, 6))
sns.barplot(x=categories, y=values, palette='pastel')
plt.ylabel('Accuracy')
plt.title('SVM Classifier Accuracy')
plt.ylim([0, 1]) # Set y-axis limits to match accuracy range (0 to 1)
plt.show()
```

RF Testing Accuracy: 0.756683033278778

RF Training Accuracy: 0.7820985127575385




Naive Bayes Classification and Accuracy Visualization

A Multinomial Naive Bayes (NB) classifier with a specified alpha value is applied to the training data. It calculates and displays both testing and training accuracies for the NB model. The code also visualizes the accuracies using a bar chart, illustrating the performance of the NB model on the dataset.

```
In [46]: nb_classifier = OneVsRestClassifier(MultinomialNB())
nb_classifier.fit(X_train, y_train)
nb_preds = nb_classifier.predict(X_test)
nb_accuracy = accuracy_score(y_test, nb_preds)
nb_train_preds = nb_classifier.predict(X_train)
nb_train_accuracy = accuracy_score(y_train, nb_train_preds)
print(f"NB Testing Accuracy: {nb_accuracy}")
print(f"NB Training Accuracy: {nb_train_accuracy}")
categories = ['Testing Accuracy', 'Training Accuracy']
values = [nb_accuracy, nb_train_accuracy]
plt.figure(figsize=(8, 6))
sns.barplot(x=categories, y=values, palette='pastel')
plt.ylabel('Accuracy')
plt.title('SVM Classifier Accuracy')
plt.ylim([0, 1]) # Set y-axis limits to match accuracy range (0 to 1)
plt.show()
```

NB Testing Accuracy: 0.7135842880523732

NB Training Accuracy: 0.7228817028243962



Neural Network Classification and Accuracy Visualization

A Multi-Layer Perceptron (MLP) classifier with specific hidden layer sizes and alpha value is applied to the training data. It calculates and displays both testing and training accuracies for the MLP model. Additionally, the code visualizes the accuracies using a bar chart, illustrating the performance of the Neural Network model on the dataset.

```
In [47]: nn_classifier = OneVsRestClassifier(MLPClassifier(hidden_layer_sizes=(64,32)))
nn_classifier.fit(X_train, y_train)
nn_preds = nn_classifier.predict(X_test)
nn_accuracy = accuracy_score(y_test, nn_preds)
nn_train_preds = nn_classifier.predict(X_train)
nn_train_accuracy = accuracy_score(y_train, nn_train_preds)
print(f"NN Testing Accuracy: {nn_accuracy}")
print(f"NN Training Accuracy: {nn_train_accuracy}")
categories = ['Testing Accuracy', 'Training Accuracy']
values = [nn_accuracy, nn_train_accuracy]
plt.figure(figsize=(8, 6))
sns.barplot(x=categories, y=values, palette='pastel')
plt.ylabel('Accuracy')
plt.title('SVM Classifier Accuracy')
plt.ylim([0, 1]) # Set y-axis limits to match accuracy range (0 to 1)
plt.show()
```

NN Testing Accuracy: 0.8041462084015275

NN Training Accuracy: 0.8706508391322145

